

Homework 2: DDPG and TRPO

Submission Guidelines: Your deliverables shall consist of 2 separate files – (i) A PDF file: Please compile all your write-ups into one .pdf file (photos/scanned copies are acceptable; please make sure that the electronic files are of good quality and reader-friendly); (ii) A zip file: Please compress all your source code into one .zip file. Please submit your deliverables via E3.

Problem 1 (Surrogate Function in TRPO)

(15 points)

In this problem, we will prove some key properties of the surrogate function used in TRPO. Recall from the slides of Lecture 16: Assume that both the state space and the action space are finite. The surrogate function $L_{\pi_{\theta_1}}(\pi_\theta)$ is defined as

$$L_{\pi_{\theta_1}}(\pi_\theta) := \eta(\pi_{\theta_1}) + \sum_{s \in \mathcal{S}} d_{\mu}^{\pi_{\theta_1}}(s) \sum_{a \in \mathcal{A}} \pi_\theta(a|s) A^{\pi_{\theta_1}}(s, a). \quad (1)$$

Show that $L_{\pi_{\theta_1}}(\pi_\theta)$ satisfies the two properties: (i) $L_{\pi_{\theta_1}}(\pi_{\theta_1}) = \eta(\pi_{\theta_1})$ and (ii) $\nabla_\theta L_{\pi_{\theta_1}}(\pi_\theta)|_{\theta=\theta_1} = \nabla_\theta \eta(\pi_\theta)|_{\theta=\theta_1}$.

Problem 2 (Solving TRPO Under Approximation Using Duality)

(20+15=35 points)

Recall from the slides of Lecture 16: We would like to solve the following optimization problem (denoted by (OPT)):

$$\text{Minimize}_{\theta \in \mathbb{R}^d} \quad -(\nabla_\theta L_{\theta_k}(\theta)|_{\theta=\theta_k})^T(\theta - \theta_k) \quad (2)$$

$$\text{subject to} \quad \frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) - \delta \leq 0. \quad (3)$$

We use θ^* to denote an optimizer of the above *primal* optimization problem (2)-(3). Note that in the above we write “Minimize” instead of “Maximize” simply to follow the conventions of the literature of optimization theory. Here we focus on the case where H is a *positive definite* matrix to avoid the technicalities (while H is only *non-negative definite* in general).

Based on the optimization theory, (OPT) is a convex optimization problem as both the objective and the constraints are convex functions. In this case, one standard way is to convert the constrained (OPT) into an unconstrained *dual* problem by defining the *Lagrangian* $\mathcal{L}(\theta, \lambda)$ and the *dual function* $D(\lambda)$ as:

$$\mathcal{L}(\theta, \lambda) := -(\nabla_\theta L_{\theta_k}(\theta)|_{\theta=\theta_k})^T(\theta - \theta_k) + \lambda \left(\frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) - \delta \right) \quad (4)$$

$$D(\lambda) := \min_{\theta \in \mathbb{R}^d} \mathcal{L}(\theta, \lambda), \quad (5)$$

where λ is called the *Lagrange multiplier*. Moreover, we call the following the *dual problem* of (OPT):

$$\max_{\lambda \geq 0} D(\lambda). \quad (6)$$

For ease of notation, we define $\lambda^* := \arg \max_{\lambda \geq 0} D(\lambda)$ as the optimizer of the dual problem. By the standard theory of convex optimization, if there exists one strictly feasible point in (OPT), then the optimal value of the dual problem is equal to the original problem (usually called “strong duality”). Moreover, if strong duality holds and a dual optimal solution λ^* exists, then any optimizer of the primal problem is also a minimizer of $\mathcal{L}(\theta, \lambda^*)$, i.e., $\theta^* = \arg \min_{\theta} \mathcal{L}(\theta, \lambda^*)$. For more details on duality, please refer to Chapter 5 of https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf.

(a) In this problem, please show that the dual function $D(\lambda)$ of (OPT) can be written as:

$$D(\lambda) = \frac{-1}{2\lambda} ((\nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k})^T H^{-1} (\nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k})) - \lambda\delta. \quad (7)$$

Accordingly, please find out λ^* based on (7).

(b) By the λ^* found in (a) and the property $\theta^* = \arg \min_{\theta} \mathcal{L}(\theta, \lambda^*)$, show that $\theta^* = \theta_k + \alpha H^{-1} \nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k}$. What is the step size α ?

Problem 3 (Deep Deterministic Policy Gradient for Continuous Control) (30+20=50 points)

In this problem, we will implement the deep deterministic policy gradient (DDPG) algorithm with the help of neural function approximators: You may write your code in either PyTorch or TensorFlow (though the sample code presumes PyTorch framework). Moreover, you are recommended to use Tensorboard to keep track of the loss terms and other related quantities of your implementation. If you are a beginner in learning the deep learning framework, please refer to the following tutorials:

- PyTorch: <https://pytorch.org/tutorials/>
- Tensorflow: <https://www.tensorflow.org/tutorials>

For the deliverables, please submit the following:

- Technical report: Please summarize all your experimental results in 1 single report (and please be brief)
- All your source code
- Your well-trained models (both the actor and critic networks) saved in either .pth files or .ckpt files

(a) We start by solving the simple “Pendulum-v0” problem (<https://gym.openai.com/envs/Pendulum-v0/>) using the DDPG algorithm. Read through `ddpg.py` and then implement the member functions of the classes **Actor**, **Critic**, and **DDPG** as well as the function **train**. Please briefly summarize your results (including the snapshots of the Tensorboard record) in the report and document all the hyperparameters (e.g. learning rates, NN architecture, and batch size) of your experiments. (Note: Pendulum is a rather basic environment mostly for the purpose of sanity check. As a result, typically it would take no more than 300 episodes to reach a well-performing policy)

(b) Based on the code for (a), adapt your DDPG algorithm to solve the “LunarLanderContinuous-v2” problem (<https://gym.openai.com/envs/LunarLanderContinuous-v2/>). Save your code in another file named `ddpg_lunarlander.py`. Please add comments to your code whenever needed for better readability. Again, please briefly summarize your results in the report and document all the hyperparameters of your experiments. (Note: As LunarLander is a more challenging environment, it would take a bit more training time to reach a well-performing policy, say about 1000 episodes or so. Also, it might require some efforts to tune the hyperparameters, e.g., learning rates and batch size. You could either do grid search or even use some more advanced techniques such as the evolutionary methods. As the main purpose of this homework is to help you get familiar with RL implementation for continuous control, there is NO hard requirement on the obtained returns of this LunarLander task. Just try your best and enjoy!)